Final Submission Audit

Status and evidence summary | 15 October 2026

Mandatory criterion	Status	File / test evidence
1. Core app and request flow	**COMPLETE**	Frontend/backend request flow and dashboard are present. Evidence: frontend/, backend/app/main.py; local health checks returned HTTP 200.
2. Routing and decision logic	**COMPLETE**	Confidence and complexity routing implemented. Evidence: backend routing modules; covered by tests/test_task6_coverage.py.
3. Model provider integration	**COMPLETE**	Provider call and error handling are implemented. Evidence: backend provider modules; mocked provider coverage in tests/test_task6_coverage.py.
4. Cost tracking	**COMPLETE**	Token usage, configured pricing, cost/savings, and unavailable/zero-baseline handling documented and tested. Evidence: backend cost modules; tests/test_task6_coverage.py.
5. Evaluation	**PARTIAL**	40-example evaluation results exist, but one router request failed (extraction_05), and evaluation was not rerun in the final audit to avoid paid calls. Evidence: data/evaluation_results.csv; README.md.
6. Automated tests	**COMPLETE**	Offline suite: 36 passed. Evidence: tests/run_offline_tests.py, tests/test_task6_coverage.py.
7. Security and Git audit	**COMPLETE**	Environment files ignored, example config retained; tracked files and reachable Git history scanned without finding credential signatures. No history rewrite or push.
8. README and deployment docs	**PARTIAL**	Commands, architecture, evaluation method, pricing assumptions, and Render setup documented. Actual live Render settings/logs unavailable, so localhost:5432 source could not be externally confirmed. Evidence: README.md.
9. Final submission audit	**COMPLETE**	Criteria mapped with evidence; backend and frontend health checked; tests and evaluation artifacts reviewed. GitHub push/history rewrite require separate authorization.

Evaluation: 40 examples; Router 71.32% accuracy, $0.00805373 cost, 39 successful / 1 failed. Tests: 36 passed.